

AI in Mathematics

Lecture 11

Reinforcement Learning

Bar-Ilan University
Nebius Academy | Stevens Institute of
Technology
June 10, 2025

About This Course

~~1 week: Intro~~


~~2 weeks: Classic ML~~

~~2 weeks: Deep Learning in Mathematics~~

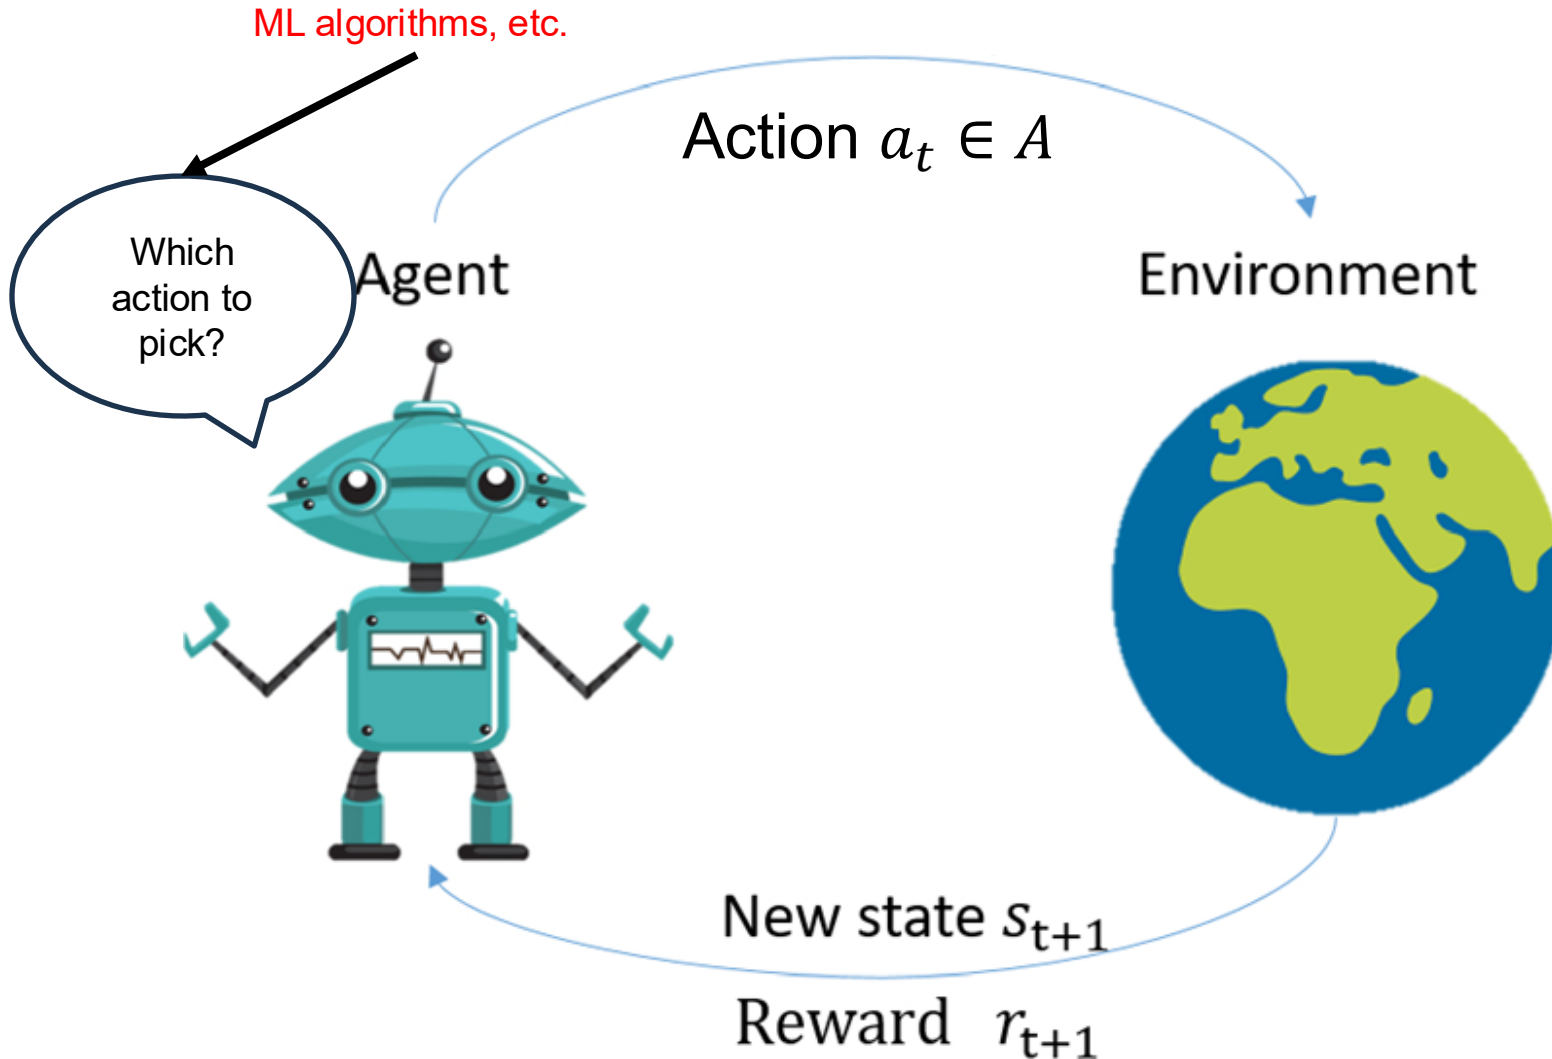
~~4 weeks: Math as an NLP problem (LLMs etc.)~~

4 weeks: Reinforcement Learning (RL) in Math

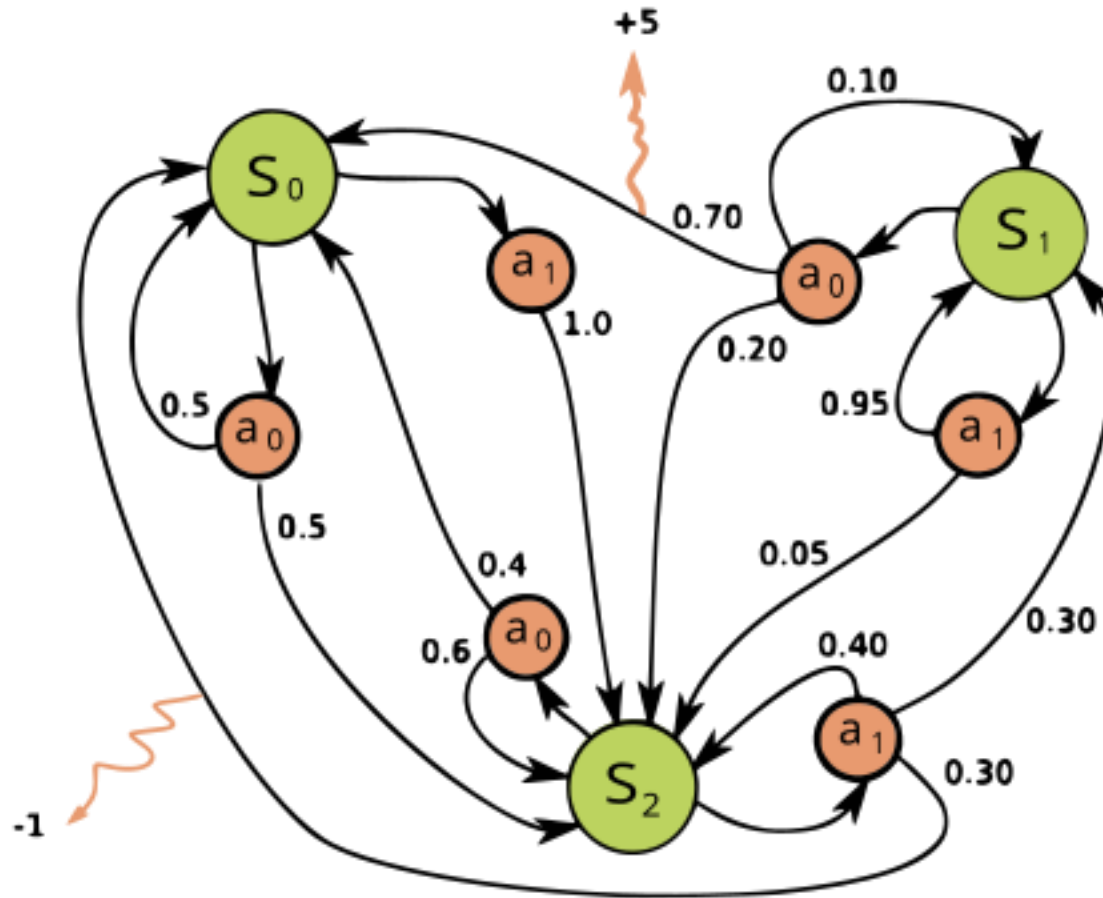
Reinforcement Learning



Reinforcement Learning

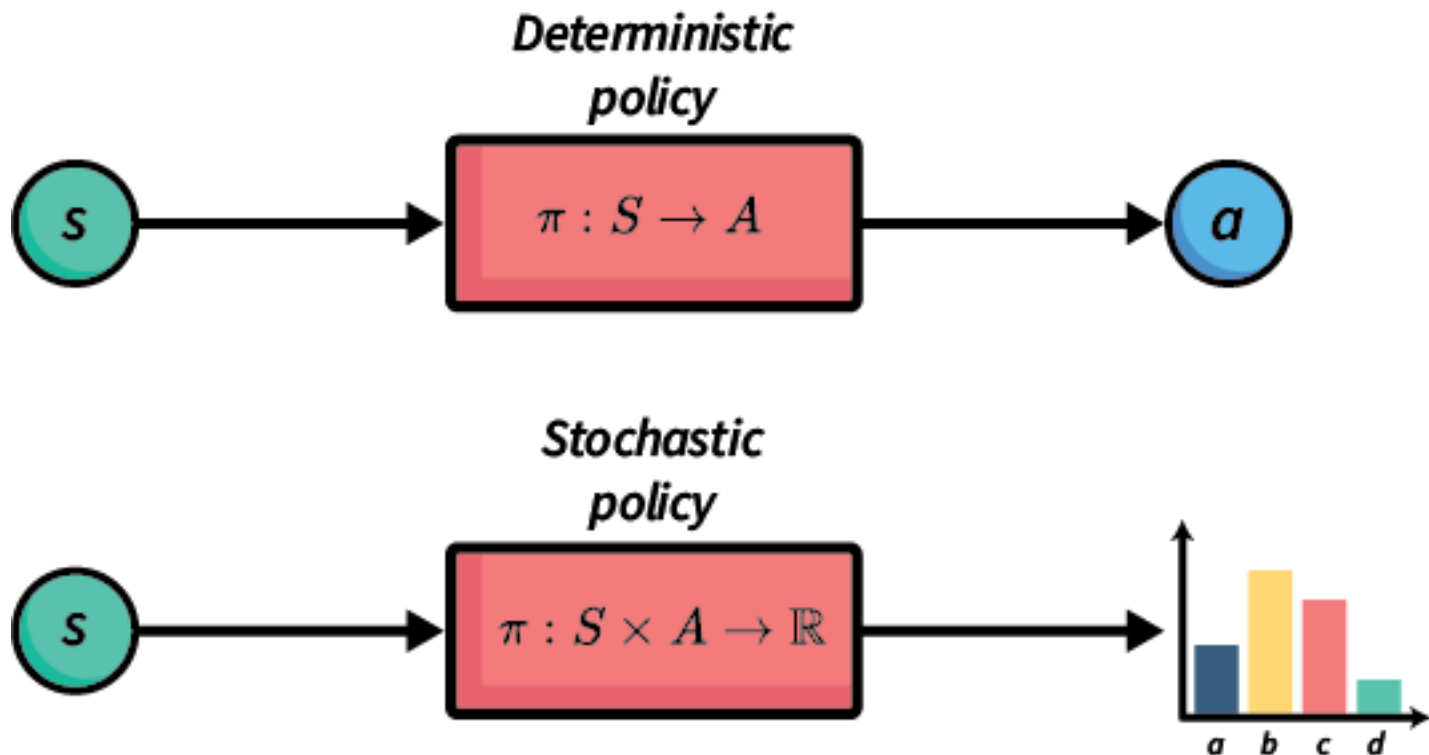


Markov Decision Process



Policy

A **policy** is the agent's strategy for choosing actions.



Policy

A **policy** is the agent's strategy for choosing actions.

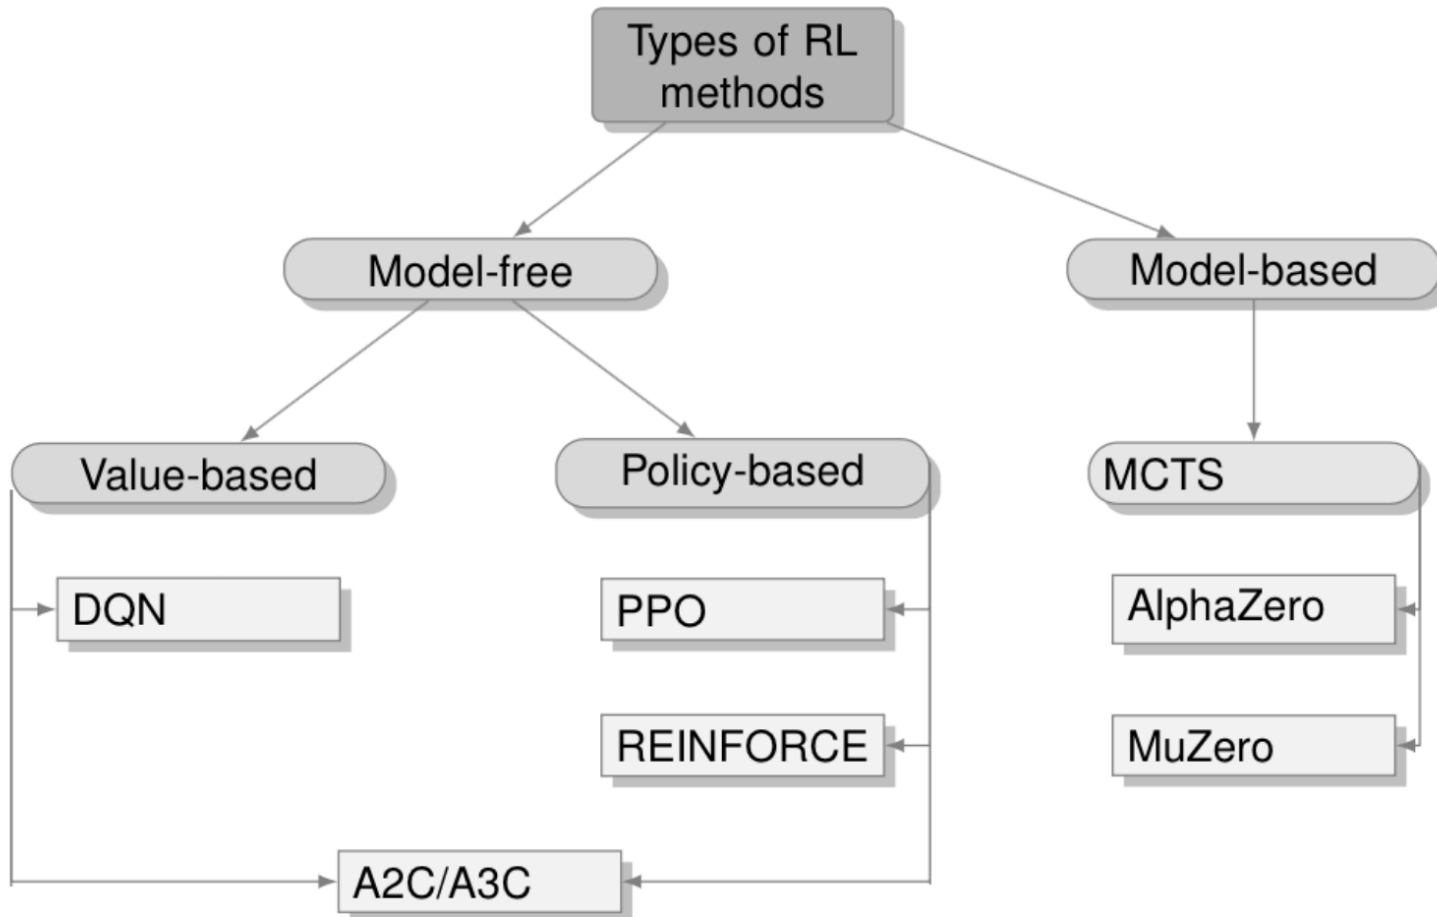
Deterministic policy:

$\pi(s) = a$ — always choose action a in state s .

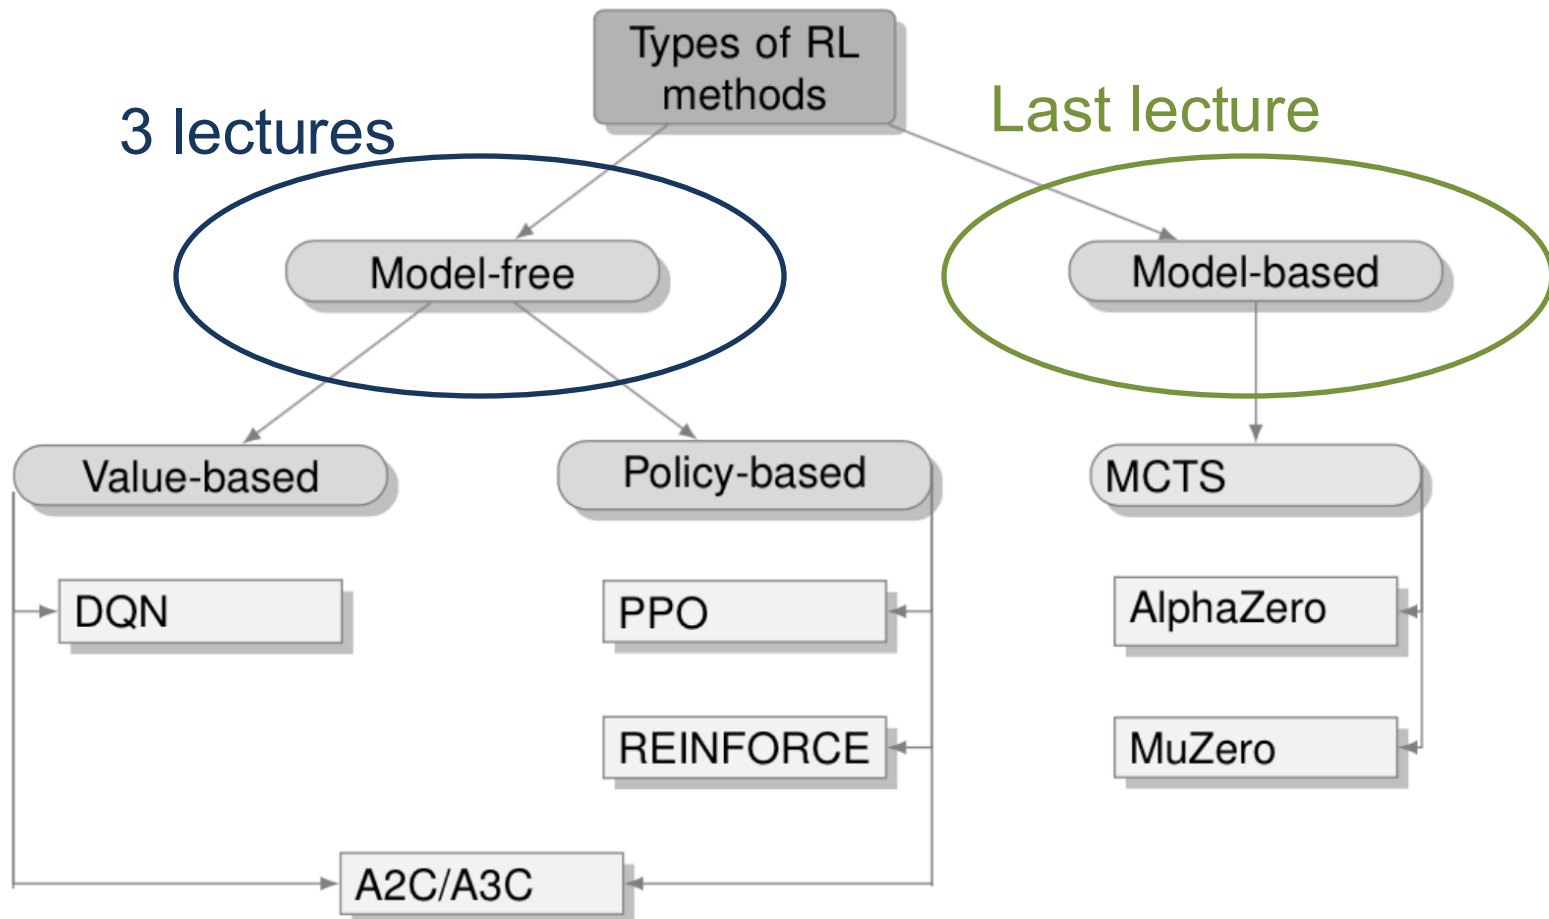
Stochastic policy:

$\pi(a | s) = P(a_t = a | s_t = s)$ — the probability of taking an action a in state s .

RL Algorithms



RL Algorithms



Best Policy

The best policy looks as follows:

$$\pi^* = \operatorname{argmax}_{\pi} \mathbb{E}_{\tau \sim \pi} \left[\sum_{t=0}^n R(s_t, a_t) \right],$$

$$a_t = \pi(s_t)$$

What are the problems with this formula?

Best Policy

$$\pi^* = \operatorname{argmax}_{\pi} \mathbb{E}_{\tau \sim \pi} \left[\sum_{t=0}^{\infty} R(s_t, a_t) \right], a_t = \pi(s_t)$$

$R(s_t, a_t)$ is a notation for $\mathbb{E}_{s_{t+1}}[R(s_t, a_t, s_{t+1})]$

$\mathbb{E}[\sum_{t=0}^{\infty} R(s_t, a_t)]$ can be infinite.

Methods to Penalize Agent

Per-step penalty (negative reward per time step)

Time limit truncation (implicit penalty)

Discount factor trick ($\gamma < 1$).

Best Policy

$$\pi^* = \operatorname{argmax}_{\pi} \mathbb{E}_{\tau \sim \pi} \left[\sum_{t=0}^{\infty} \gamma^t R(s_t, a_t) \right],$$
$$a_t = \pi(s_t)$$

$$\mathcal{J}(\pi_{\theta}) = \mathbb{E}_{\tau \sim \pi_{\theta}} [\sum_{t=0}^{\infty} \gamma^t R(s_t, a_t)].$$

So we can apply gradient descent on $\mathcal{J}(\pi_{\theta})$ to find π^* .

Policy Gradient

$$\theta_{k+1} = \theta_k + \alpha \nabla_{\theta} \mathcal{J}(\pi_{\theta_k})$$

We need to calculate $\nabla_{\theta} \mathcal{J}(\pi_{\theta})$ first.

$$\begin{aligned} \mathcal{J}(\pi_{\theta}) &= \mathbb{E}_{\tau \sim \pi_{\theta}} \left[\sum_{t=0}^{\infty} \gamma^t R(s_t, a_t) \right] = \mathbb{E}_{\tau \sim \pi_{\theta}} [R(\tau)] \\ &= \int R(\tau) p_{\theta}(\tau) d\tau \end{aligned}$$

We need to understand how to compute the probability $p_{\theta}(\tau)$.

Policy Gradient

$$J(\pi_\theta) = \mathbb{E} \left[\sum_{t=0}^{\infty} \gamma^t R(s_t, a_t) \right] = \int R(\tau) p_\theta(\tau) d\tau$$

$$p_\theta(\tau) = Pr(s_0) \cdot \prod_{t=0}^{\infty} \pi_\theta(a_t | s_t) \cdot Pr(s_{t+1} | s_t, a_t)$$

So we can not simply calculate it or its gradient.
We need to invent some other way.

Policy Gradient

$$\nabla_{\theta} \mathcal{J}(\pi_{\theta}) = \int R(\tau) \nabla_{\theta} p_{\theta}(\tau) d\tau$$

$$= \int R(\tau) \nabla_{\theta} p_{\theta}(\tau) \frac{p_{\theta}(\tau)}{p_{\theta}(\tau)} d\tau$$

$$= \int R(\tau) \nabla_{\theta} \ln p_{\theta}(\tau) p_{\theta}(\tau) d\tau$$

$$= \mathbb{E}_{\tau} [R(\tau) \nabla_{\theta} \ln p_{\theta}(\tau)]$$

Policy Gradient

Why is it better?

$$\begin{aligned}\nabla_{\theta} \ln p_{\theta}(\tau) &= \\&= \nabla_{\theta} \ln \left[p(s_0) \cdot \prod_{t=0}^{\infty} \pi_{\theta}(a_t | s_t) p(s_{t+1} | s_t, a_t) \right] \\&= \nabla_{\theta} \sum_{t=0}^{\infty} \ln \pi_{\theta}(a_t | s_t)\end{aligned}$$

Policy Gradient

$$\begin{aligned}\nabla_{\theta} \mathcal{J}(\pi_{\theta}) &= \mathbb{E}_{\tau} [R(\tau) \nabla_{\theta} \ln p_{\theta}(\tau)] = \\ &= \mathbb{E}_{\tau} \left[R(\tau) \nabla_{\theta} \sum_{t=0}^{\infty} \ln \pi_{\theta}(a_t | s_t) \right] \\ &= \mathbb{E}_{\tau} \left[\sum_t \nabla_{\theta} \ln \pi_{\theta}(a_t | s_t) \sum_{t' \geq t} \gamma^{t'-t} R(s_{t'}, a_{t'}) \right] \\ Q(s_t, a_t) &= \sum_{t' \geq t} \gamma^{t'-t} R(s_{t'}, a_{t'})\end{aligned}$$

Policy Gradient

$$\begin{aligned}\nabla_{\theta} \mathcal{J}(\pi_{\theta}) &= \mathbb{E}_{\tau} \left[\sum_t \nabla_{\theta} \ln \pi_{\theta}(a_t | s_t) Q(s_t, a_t) \right] \\ &= \mathbb{E}_{(s, a) \sim d^{\pi_{\theta}}(s) \pi_{\theta}(s | a)} [\nabla_{\theta} \ln \pi_{\theta}(a | s) Q(s, a)]\end{aligned}$$

Where the last expectation is over the stationary distribution of states defined by our policy.

Policy Gradient

$$\nabla_{\theta} \mathcal{J}(\pi_{\theta}) = \mathbb{E}_{(s, a) \sim d^{\pi_{\theta}}(s) \pi_{\theta}(a | s)} [\nabla_{\theta} \ln \pi_{\theta}(a | s) Q(s, a)]$$

This is the policy gradient theorem.

It expresses the gradient of expected return with respect to the policy parameters as an expectation over state-action pairs.

REINFORCE

$$\nabla_{\theta} \mathcal{J}(\pi_{\theta}) = \mathbb{E}_{\pi_{\theta}} \left[\nabla_{\theta_k} \ln \pi_{\theta_k}(a|s) Q(s, a) \right]$$

In practice, REINFORCE algorithm approximates this gradient by trajectories sampled from the policy.

We get $Q(s, a)$ from a trajectory we have seen.

REINFORCE

$$\nabla_{\theta} \mathcal{J}(\pi_{\theta}) = \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta_k} \ln \pi_{\theta_k}(a|s) Q(s, a)]$$

We need full trajectories to be able to compute $Q(s, a)$.

This method suffers from high variance, as the gradient estimate $\nabla_{\theta} \mathcal{J}(\pi_{\theta})$ may fluctuate substantially across different episodes.

Variance problem

$$\nabla_{\theta} \mathcal{J}(\pi_{\theta}) = \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \ln \pi_{\theta}(a|s) (Q(s, a) - b(s))]$$

$b(s)$ is some function of the state that allows us to reduce variance.

One good solution is to use $b(s) = V^{\pi_{\theta}}(s) = \mathbb{E}_a Q(s, a)$. It is exactly what A2C will do.

So $V^{\pi_{\theta}}(s)$ represents how good or promising this state is on average, under the future path defined by π .

A2C

$$\nabla_{\theta} \mathcal{J}(\pi_{\theta}) \approx \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \ln \pi_{\theta}(a|s) A(s, a)]$$

where the *advantage* function is defined as

$$A(s, a) = Q(s, a) - V^{\pi_{\theta}}(s), \quad V^{\pi_{\theta}}(s) = \mathbb{E}_a Q(s, a).$$

To approximate A , we will use

$$A_t(s, a) = r_t + \gamma V(s_{t+1}) - V(s_t)$$

This is a **one-step temporal difference (TD) estimate** of the advantage — efficient and requires minimal storage.

1. The **value function** $V(s)$ is learned using a neural network.
2. The **advantage estimate** A_t uses only the **immediate reward** and **next state value** — no need to store full trajectories.

A2C

Collect Trajectories

1. Interact with the environment using the current policy π_θ .
2. Collect (s_t, a_t, r_t, s_{t+1}) for multiple time steps (usually in parallel environments).

Use a neural network (critic) to predict $V_\phi(s) \approx V^\pi(s)$.

Use 1-step TD estimate:

$$A_t = r_t + \gamma V(s_{t+1}) - V(s_t)$$

Update Policy (Actor):

$$1. \theta \leftarrow \theta + \alpha \cdot \nabla_\theta \ln \pi_\theta(a_t | s_t) \cdot A_t$$

Update Critic using loss function:

$$L_{critic} = (r_t + \gamma V(s_{t+1}) - V(s_t))^2$$

A2C (Detailed)

Collect Trajectories (Shared Step)

- Interact with the environment using the current policy π_θ .
- Collect transitions (s_t, a_t, r_t, s_{t+1}) over multiple steps (e.g., from parallel environments)
- Use a neural network (the critic) to predict $V_\phi(s) \approx V^\pi(s)$.
- Compute 1-step TD Advantage estimate:

$$A_t = r_t + \gamma V(s_{t+1}) - V(s_t)$$

What Actually Happens:

We **collect a batch** of transitions (s_t, a_t, r_t, s_{t+1}) from multiple environments or steps.

Then compute:

- **Batch of value predictions** $V(s_t), V(s_{t+1})$
- **Batch of advantages** $A_t = r_t + \gamma V(s_{t+1}) - V(s_t)$

These batches are then used to perform **gradient updates**

A2C (Detailed)

Actor Update

Policy Improvement

The actor updates the policy parameters θ to maximize expected return.

Gradient ascent step:

$$\theta \leftarrow \theta + \alpha \cdot \nabla_{\theta} \ln \pi_{\theta}(a_t | s_t) \cdot A_t$$

Critic Update

The critic learns to estimate the **state value** $V^{\pi}(s) \approx V_{\phi}(s)$

Loss function:

$$L_{critic} = \left(r_t + \gamma V_{\phi}(s_{t+1}) - V_{\phi}(s_t) \right)^2$$

Gradient descent step:

$$\phi \leftarrow \phi - \eta \cdot \nabla_{\phi} L_{critic}$$

In both cases we can use stochastic gradient descent to update parameters.